

hw3

October 20, 2024

1 Problem 4

```
[117]: from sklearn.datasets import load_digits
from sklearn.metrics import log_loss, accuracy_score
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import GridSearchCV

import matplotlib.pyplot as plt
```

```
[85]: digits = load_digits()
X = digits.data
y = digits.target

X_train, X_test = X[:1000], X[1000:]
y_train, y_test = y[:1000], y[1000:]
```

```
[ ]: deviance_dct = {
    (0.1, 0.5): None,
    (0.1, 1) : None,
    (0.01, 0.5): None,
    (0.01, 1): None
}

for lr in [0.1, 0.01]:
    for p in [0.5, 1]:
        gb_clf = GradientBoostingClassifier(n_estimators=1500, # T = 1500
                                            learning_rate=lr, # 0.1 and 0.01
                                            max_leaf_nodes=4,
                                            min_samples_split=5,
                                            subsample=p, # 0.5 and 1
                                            random_state=42)

        gb_clf.fit(X_train, y_train)

        y_pred_proba_stages = gb_clf.staged_predict_proba(X_test)
```

```

deviances = []
for y_pred_proba in y_pred_proba_stages:
    deviance = log_loss(y_test, y_pred_proba)
    deviances.append(deviance)

deviance_dct[(lr, p)] = deviances

```

```

[101]: for lr in [0.1, 0.01]:
        for p in [0.5, 1]:

            plt.plot(deviance_dct[(lr,p)], label=f'$\eta$={lr}, $p$={p}')


plt.title('Deviance Comparison')
plt.legend()
plt.xlabel('Iteration')
plt.ylabel('Log Loss')

```

```

[101]: Text(0, 0.5, 'Log Loss')

```



```
[103]: param_grid = {
        'n_estimators': [20, 30, 40, 60, 80, 100, 200, 500]
    }

    gb_clf = GradientBoostingClassifier(learning_rate=0.1, # Example learning rate
                                       max_leaf_nodes=4,
                                       min_samples_split=5,
                                       subsample=0.5,      # Example subsample
                                       random_state=42)

    grid_search = GridSearchCV(estimator=gb_clf,
                               param_grid=param_grid,
                               cv=10,
                               scoring='neg_log_loss',
                               n_jobs=-1)

    grid_search.fit(X_train, y_train)

    best_n_estimators = grid_search.best_params_['n_estimators']
```

```
[108]: import pandas as pd
        pd.DataFrame(grid_search.cv_results_)
```

```
[108]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	\
0	1.113934	0.024499	0.006466	0.004081	
1	1.633280	0.044832	0.003623	0.000530	
2	2.151802	0.046762	0.003758	0.000719	
3	3.099959	0.082132	0.004153	0.000775	
4	4.124060	0.086896	0.004514	0.000750	
5	5.278415	0.084862	0.005714	0.002679	
6	10.710508	0.210531	0.006764	0.000723	
7	21.778956	0.112896	0.009756	0.000930	

	param_n_estimators	params	split0_test_score	\
0	20	{'n_estimators': 20}	-0.634718	
1	30	{'n_estimators': 30}	-0.517642	
2	40	{'n_estimators': 40}	-0.462130	
3	60	{'n_estimators': 60}	-0.414317	
4	80	{'n_estimators': 80}	-0.384714	
5	100	{'n_estimators': 100}	-0.352984	
6	200	{'n_estimators': 200}	-0.326890	
7	500	{'n_estimators': 500}	-0.484325	

	split1_test_score	split2_test_score	split3_test_score	split4_test_score	\
0	-0.410750	-0.357796	-0.463038	-0.542924	
1	-0.306412	-0.261578	-0.343196	-0.406640	
2	-0.256677	-0.211052	-0.306566	-0.337621	

3	-0.200579	-0.160666	-0.253328	-0.285509
4	-0.179854	-0.145694	-0.217918	-0.251432
5	-0.166531	-0.121922	-0.189873	-0.226205
6	-0.137073	-0.098897	-0.120130	-0.223812
7	-0.171181	-0.106960	-0.103007	-0.292377

	split5_test_score	split6_test_score	split7_test_score	split8_test_score	\
0	-0.585417	-0.483604	-0.511252	-0.511213	
1	-0.407291	-0.317340	-0.397742	-0.420775	
2	-0.341101	-0.251725	-0.328141	-0.387675	
3	-0.274873	-0.167223	-0.262561	-0.316265	
4	-0.235531	-0.128219	-0.246509	-0.282631	
5	-0.204619	-0.102296	-0.232413	-0.245616	
6	-0.159400	-0.048569	-0.227744	-0.165468	
7	-0.157914	-0.016387	-0.301402	-0.141732	

	split9_test_score	mean_test_score	std_test_score	rank_test_score
0	-0.462232	-0.496294	0.076584	8
1	-0.341912	-0.372053	0.069221	7
2	-0.290532	-0.317322	0.068645	6
3	-0.212774	-0.254810	0.072020	5
4	-0.179247	-0.225175	0.070668	4
5	-0.159207	-0.200167	0.067655	3
6	-0.117145	-0.162513	0.075221	1
7	-0.160283	-0.193557	0.125934	2

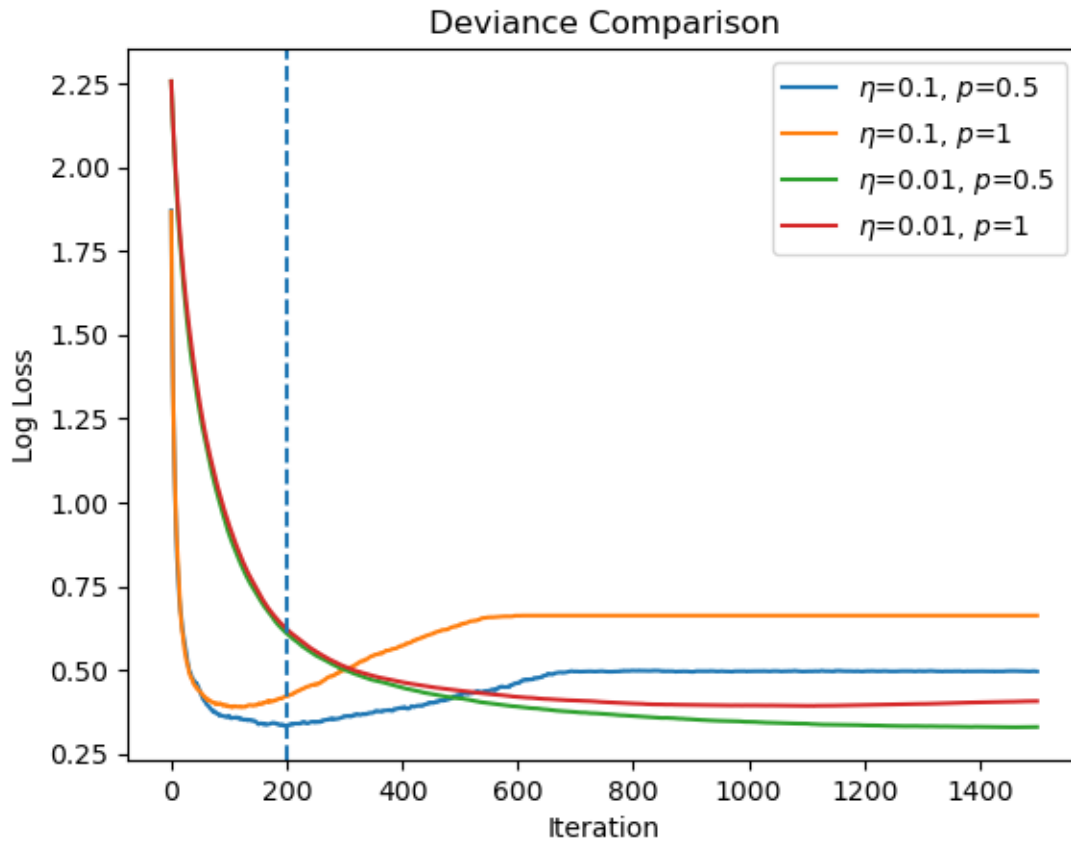
```
[115]: for lr in [0.1, 0.01]:
        for p in [0.5, 1]:

            plt.plot(deviance_dct[(lr,p)], label=f'$\eta$={lr}, $p$={p}')

plt.axvline(best_n_estimators, linestyle='--')

plt.title('Deviance Comparison')
plt.legend()
plt.xlabel('Iteration')
plt.ylabel('Log Loss')
```

```
[115]: Text(0, 0.5, 'Log Loss')
```



```
[116]: gb_clf = GradientBoostingClassifier(n_estimators=3000, # T = 1500
                                         learning_rate=0.1, # 0.1 and 0.01
                                         max_leaf_nodes=4,
                                         min_samples_split=5,
                                         subsample=1,      # 0.5 and 1
                                         random_state=42)
```

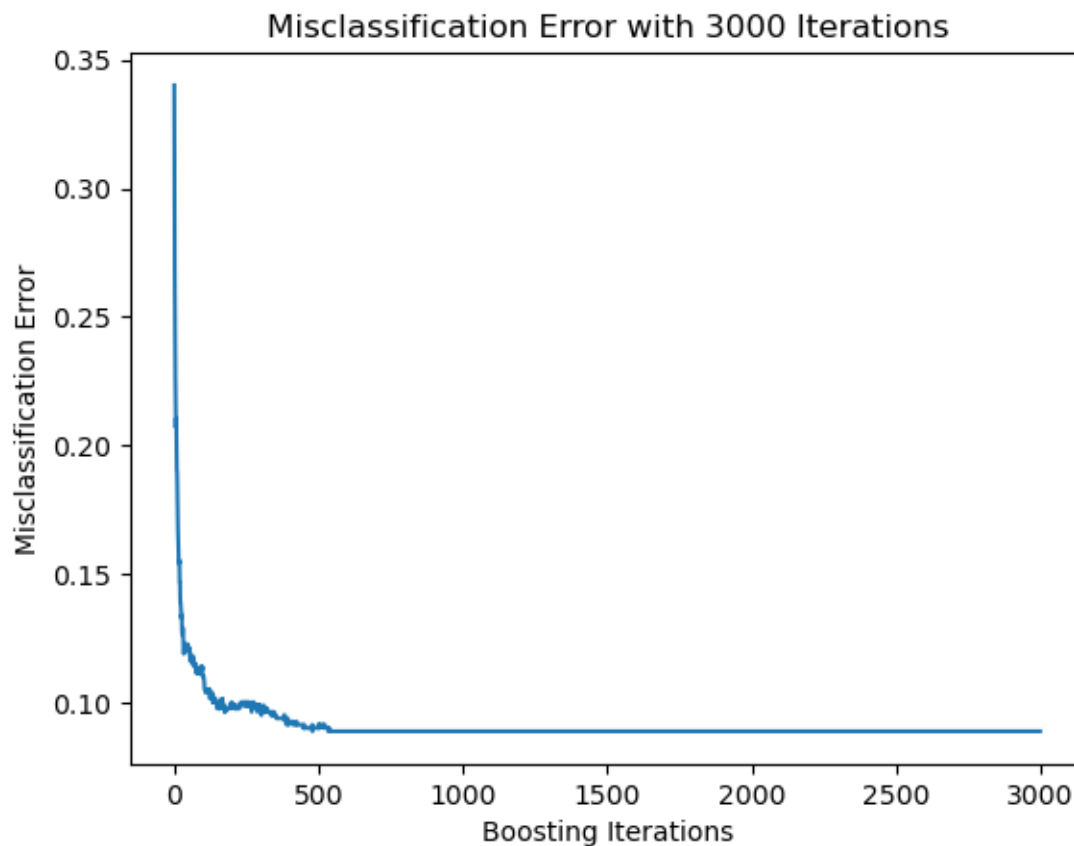
```
gb_clf.fit(X_train, y_train)
```

```
[116]: GradientBoostingClassifier(max_leaf_nodes=4, min_samples_split=5,
                                n_estimators=3000, random_state=42, subsample=1)
```

```
[124]: errors = []
for y_pred in gb_clf.staged_predict(X_test):
    errors.append(
        1 - accuracy_score(y_pred, y_test)
    )
```

```
[132]: plt.plot(errors)
# plt.yscale('log')
# plt.ylim(0,0.1)
plt.xlabel("Boosting Iterations")
plt.ylabel("Misclassification Error")
plt.title("Misclassification Error with 3000 Iterations")
```

```
[132]: Text(0.5, 1.0, 'Misclassification Error with 3000 Iterations')
```



2 Problem 5

```
[206]: from sklearn.neighbors import RadiusNeighborsClassifier
from collections import Counter
import random
import numpy as np
```

```
[142]: d=2
mu0 = [
    [-2,-3.5],
```

```

    [0,0],[2,-3.5]
]

Id = [
    [25,0],
    [0,25]
]

def Data_Generate(n):
    outX = []
    outY = np.random.choice(
        list(range(3)), n, replace=True, p=[0.5,0.25,0.25]
    )

    for i in range(n):
        mu = mu0[outY[i]]
        X = np.random.multivariate_normal(mu,Id,1)[0].tolist()
        outX.append(X)

    return np.array(outX), np.array(outY)

```

```

[183]: n = 500
testX, testY = Data_Generate(n)

```

```

[169]: def q_quantile_interpoint_distance(data, q):
    distances = np.linalg.norm(data[:, np.newaxis, :] - data[np.newaxis, :, :],
↪axis=2)
    # print(distances)
    # print(distances.shape)

    interpoint_distances = distances[np.triu_indices_from(distances, k=1)]
    # print(interpoint_distances)
    # print(interpoint_distances.shape)
    return np.quantile(interpoint_distances, q)

```

```

[ ]: NUM_EXPERIMENT = 100

well_defined_proportions = [[] for _ in range(11)]
errors = [[] for _ in range(11)]
# repeat 1000 times

for experiment in range(NUM_EXPERIMENT):
    trainX, trainY = Data_Generate(n)

    for q in range(1, 11):
        dist = q_quantile_interpoint_distance(trainX, q/10)
        # print(q/10, dist)

```

```

ENN = RadiusNeighborsClassifier(radius=dist)
ENN.fit(trainX, trainY)

neigh_indices = ENN.radius_neighbors(testX, return_distance=False)
well_defined_count = 0
for indices in (neigh_indices):
    if len(indices):
        well_defined_count += 1

# print(experiment, q, well_defined_count, well_defined_count / n)
well_defined_proportion = well_defined_count / n
well_defined_proportions[q].append(well_defined_proportion)

try:
    y_pred = ENN.predict(testX)

    test_error = 1 - accuracy_score(testY, y_pred)
    errors[q].append(test_error)

    # print(neigh_indices)
    # print(len(neigh_indices))
    # print(neigh_indices[0].shape, neigh_indices[39].shape,
    ↪neigh_indices[139].shape)
except ValueError as e:
    # print(e)
    pass

```

plot error

```

[246]: mus = []
stds = []

for q in range(1, 11):
    err = np.array(errors[q])
    mu = err.mean()
    std = err.std()

    mus.append(mu)
    stds.append(std)

x_values = np.arange(1, 11) / 10

plt.errorbar(x_values, mus, yerr=stds, fmt='o-', capsize=5)

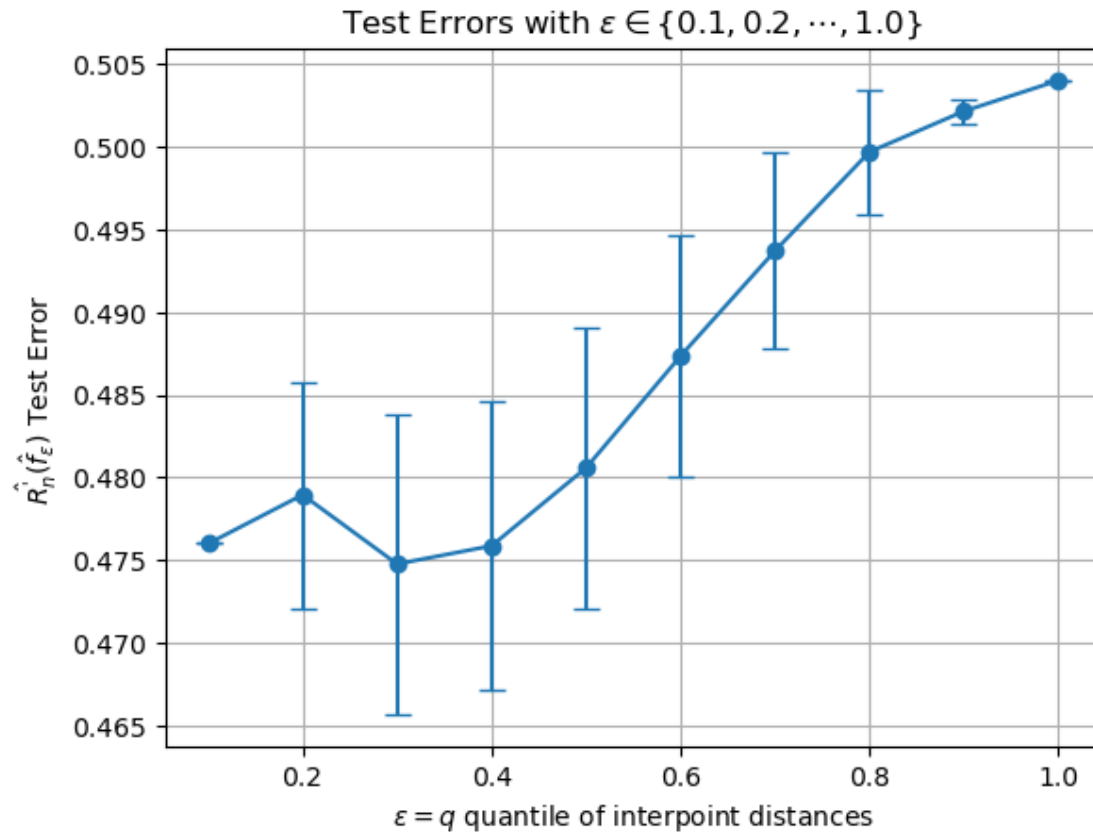
# Add labels and title
plt.xlabel("\epsilon = q$ quantile of interpoint distances")

```



```
plt.ylabel("$\hat{R}_n(\hat{f}_\epsilon)$ Test Error")
plt.title("Test Errors with $\epsilon \in \{0.1, 0.2, \dots, 1.0\}$")

# Show the plot
plt.grid(True)
plt.show()
```



```
[275]: mus = []
stds = []

for q in range(1, 11):
    arr = np.array(well_defined_proportions[q])
    # print(arr)
    mu = arr.mean()
    std = arr.std()

    mus.append(mu)
    stds.append(std)

x_values = np.arange(1, 11) / 10
```

```

plt.errorbar(x_values, mus, yerr=stds, fmt='o-', capsize=5)

# Add labels and title
plt.xlabel("$\epsilon = q$ quantile of interpoint distances")
plt.ylabel("Ratio of well-defined $\hat{f}_\epsilon$")
plt.title("Ratio of well-defined $\hat{f}_\epsilon$ with $\epsilon \in \{0.1, 0.2, \dots, 1.0\}$")

# Show the plot
plt.grid(True)
plt.show()

```

